Event collection framework @ Blinkit

Why?

We are ‘event’ driven

How?

Aggregation of ‘events’

What?

Observability


Gotta catch 'em all!
BACKEND EVENTS
FRONTEND EVENTS
imgflip.com

Frontend Events

Clicks and Impressions

Backend Events

- Across Services (Pub-Sub)
- Localised to Services (Logs)

Logs -> Aggregator -> Kafka

- Log events with our client

- Write to a file -> Stdout in a sidecar

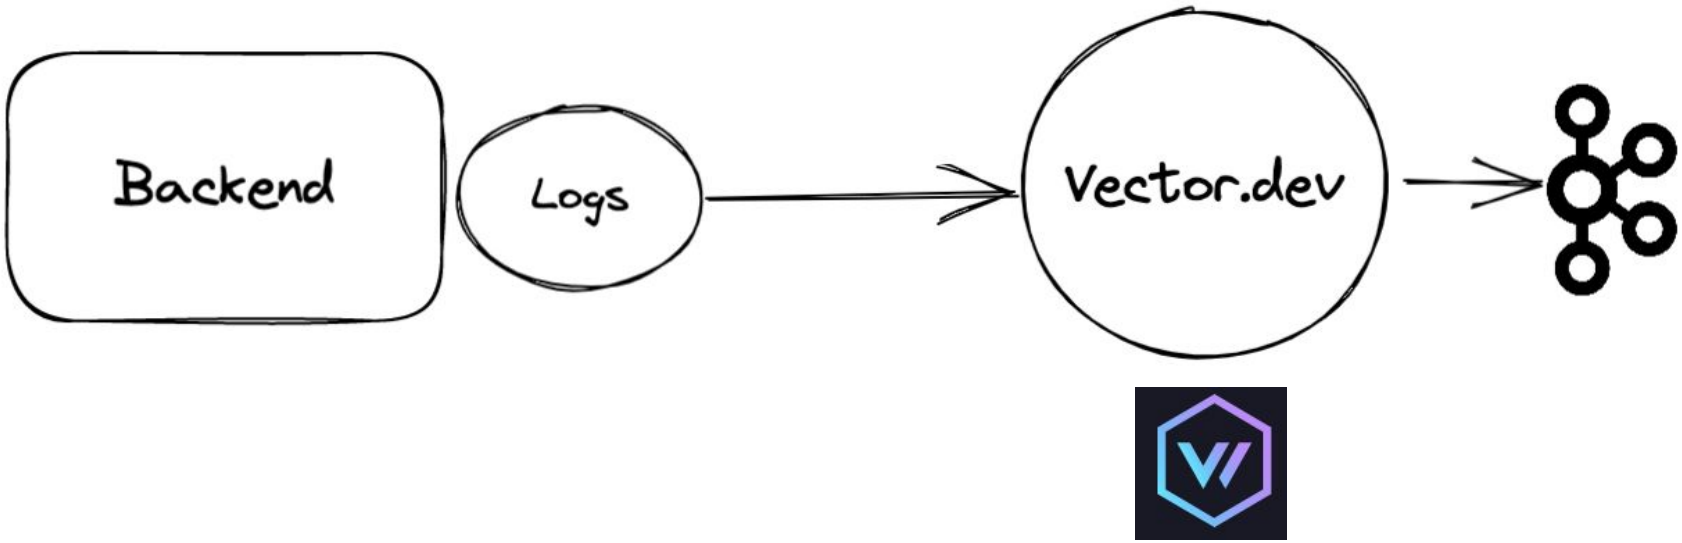
Backend
Logs
Vector.dev

Aggregation Layer

- To Pre-aggregate or Not To
- Apache Flink (Pre)
- Sink (Post)

Sink

Apache Pinot ✨

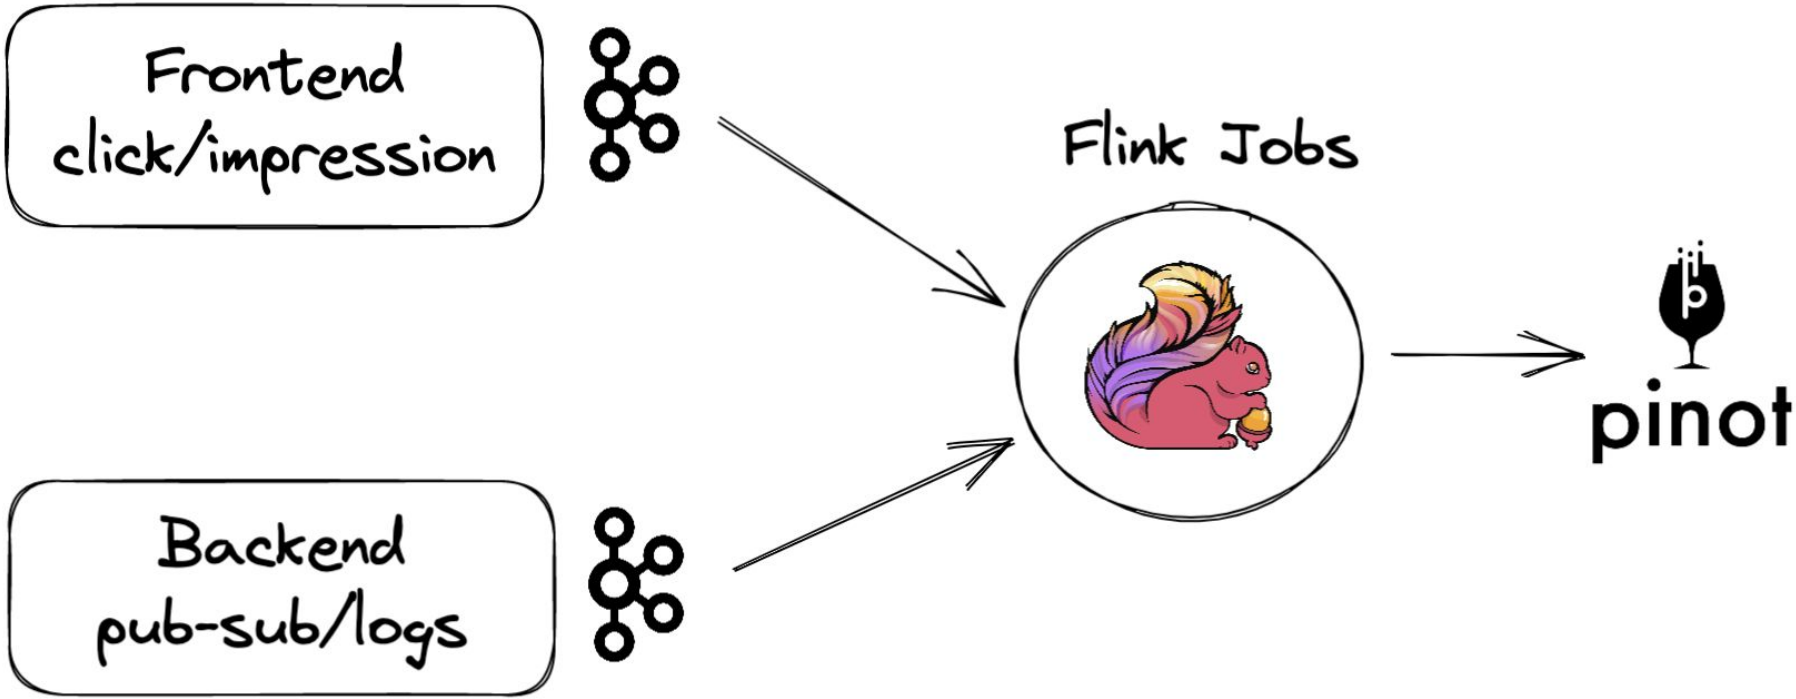
Frontend
click/impression
Backend
pub-sub/logs
Flink Jobs
pinot

Observability Overlord

Sonar

What we need

- Ingest from Kafka
- Upserts!
- Primary Key

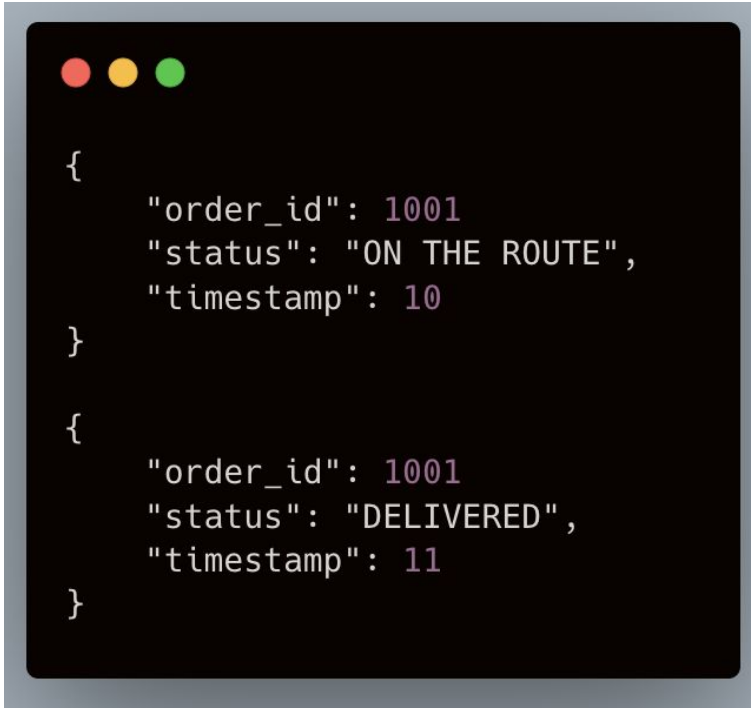

SQL-like Interface

```
SELECT
    SELECT
    SUM(total_cost) as gmv,
    COUNT(DISTINCT(cart_id)) as order_count,
    COUNT(DISTINCT customer_id) as transacting_users_count,

FROM
    order_table
WHERE
    insert_ts >= dateTrunc('DAY', now(), 'MILLISECONDS', 'Asia/Kolkata', 'MILLISECONDS')
```

Performance

- Fast Queries
- Historical Data

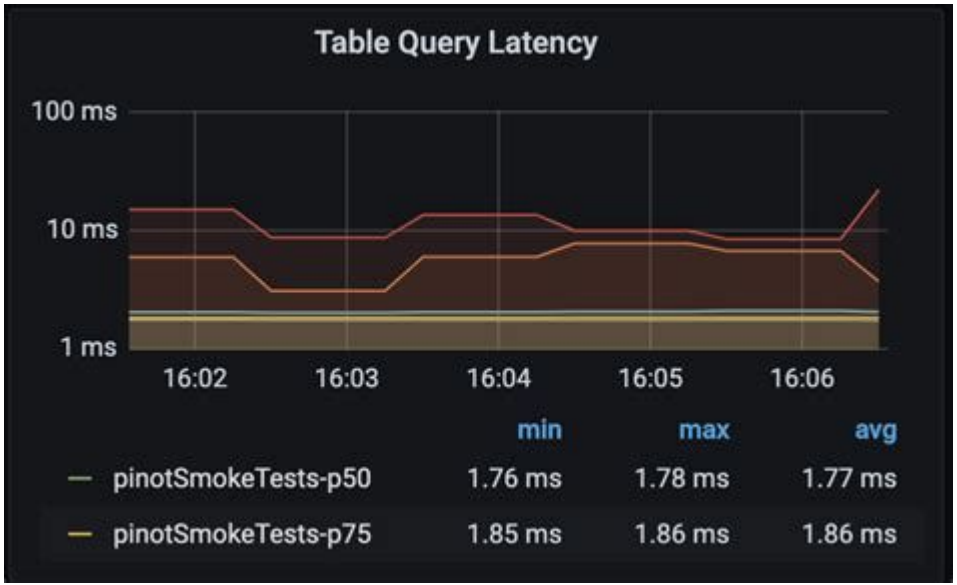

WE GOT IT!

- Pluggable indexing
- Real time ingestion
- SQL Interface
- Lookup Joins
- Hybrid tables
- Full and Partial Upserts

...

pinot

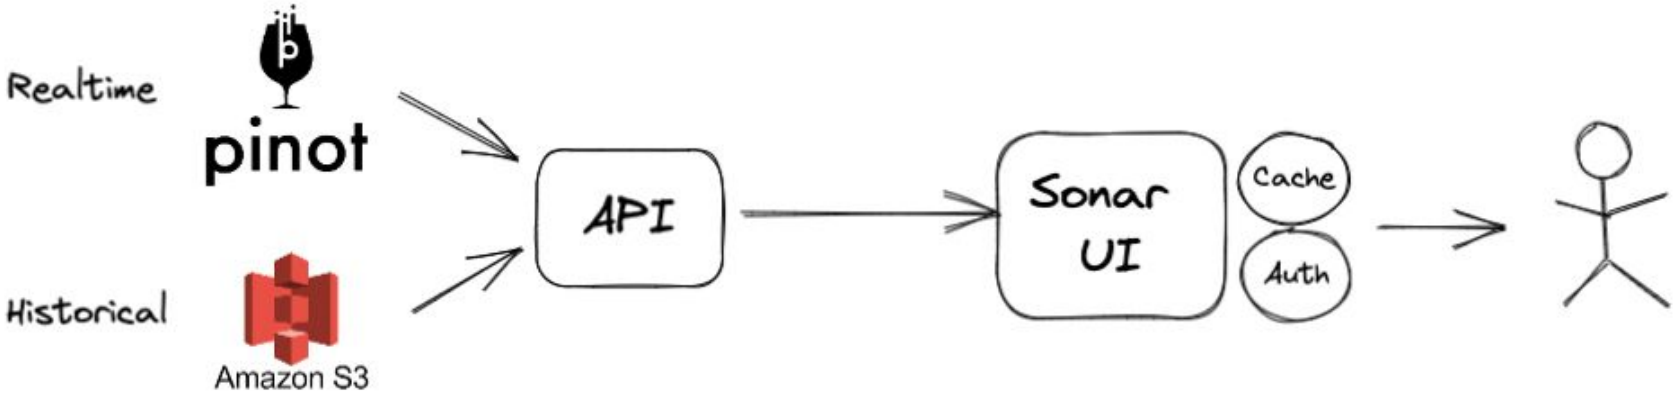
Realtime
pinot
Historical
Amazon S3
API
Sonar
UI
Cache
Auth

Ad Platform

- Performance feedback loop
- Powering campaigns
- Agg metrics for brand partners
- Rotate ads and ensure delivery rates

More where that came from...

- Performance monitoring
- App HTTP metrics
- etc...

Thank you 🙏